


Shohin assigns learned owners across inference time while freezing the host backbone

a Dense model-owned temporal lifecycle

b Sparse-host revision surface


c Tokenwise causal ownership

same source + draft
same output budget




Inference sees no correctness label, verifier, benchmark route, or tool result

final 16 layers; native router/expert trainables = 0



Mixtral: rank 18, $\alpha=18$, 3,538,944 trained parameters

Only w_i and b_i train; both trajectory residuals stay frozen



Qwen3.6-35B-A3B: 32,784 gate parameters; response loss only