

Shohin assigns learned owners across inference time while freezing the host backbone

a Dense model-owned temporal lifecycle

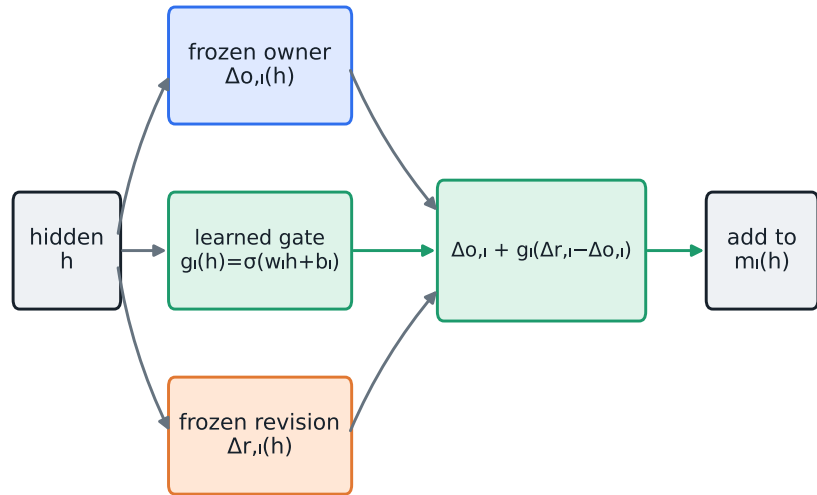
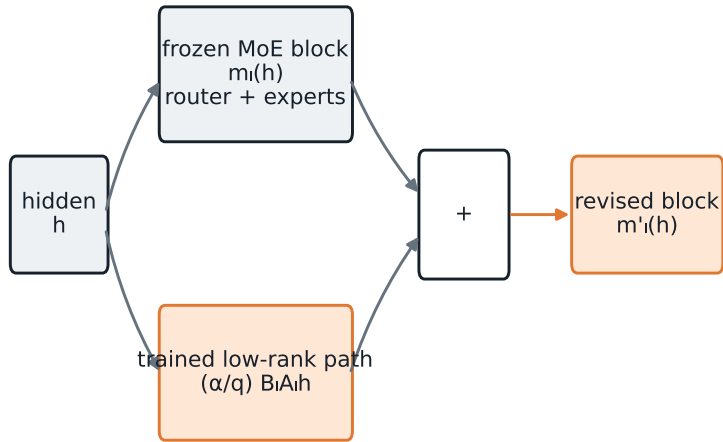
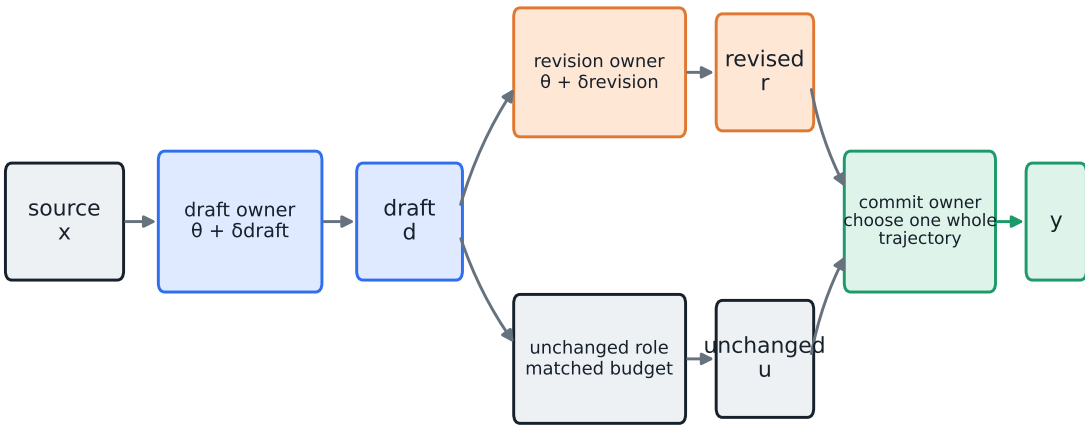
b Sparse-host revision surface

c Tokenwise causal ownership

dense lifecycle: same source + draft
same output budget

final 16 layers; native router/expert trainables = 0

Only w_i and b_i train; both trajectory residuals stay frozen



Inference sees no correctness label, verifier, benchmark route, or tool result

Mixtral: rank 18, $\alpha=18$, 3,538,944 trained parameters

Qwen3.6-35B-A3B: 32,784 incremental gate; 2,392,080 Shohin parameters total